

# SWE2001: System Program

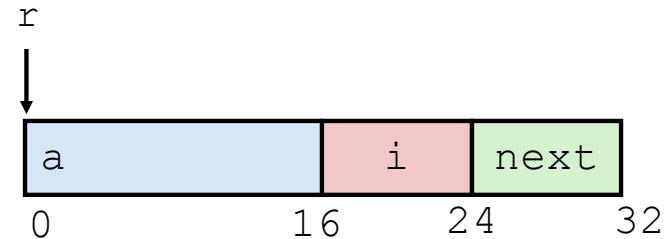
## Lecture 0x0A: Alignment

---

Systems Security Lab @ SKKU

# Structure Representation

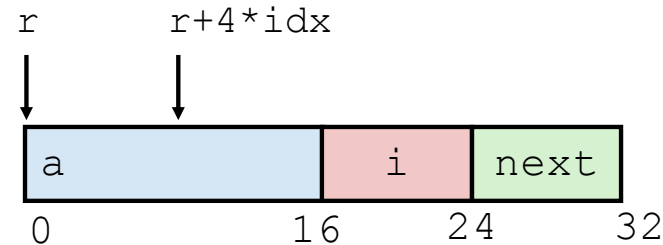
```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- ▶ Structure represented as block of memory
  - Big enough to hold all of the fields
- ▶ Fields ordered according to declaration
  - Even if another ordering could yield a more compact representation
- ▶ Compiler determines overall size + positions of fields
  - Machine-level program has no understanding of the structures in the source code

# Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



## ► Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as  $r + 4 \cdot idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```

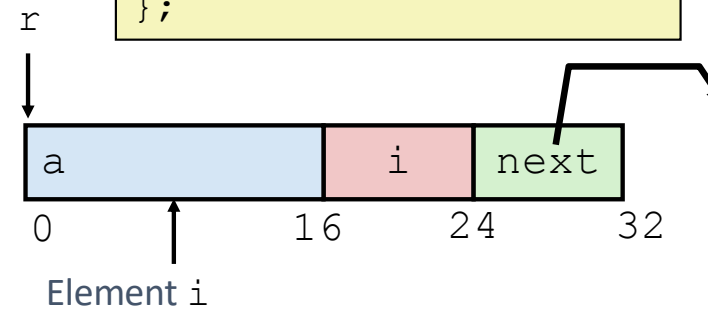
```
# r in %rdi, idx in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

# Following Linked List

## ► C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Register	Value
%rdi	r
%rsi	val

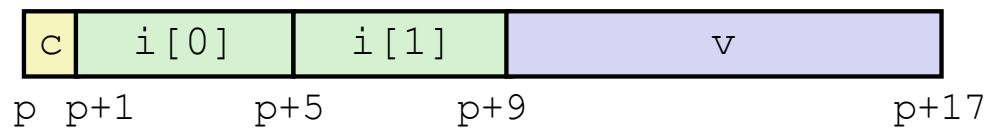
```
.L11:                                # loop:
    movslq 16(%rdi), %rax            # i = M[r+16]
    movl   %esi, (%rdi,%rax,4)       # M[r+4*i] = val
    movq   24(%rdi), %rdi           # r = M[r+24]
    testq  %rdi, %rdi               # Test r
    jne    .L11                     # if !=0 goto loop
```

# Alignment Principles

- ▶ Aligned Data
  - Primitive data type requires  $K$  bytes
  - Address must be multiple of  $K$
  - Required on some machines; advised on x86-64
- ▶ Motivation for Aligning Data
  - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
    - Inefficient to load or store datum that spans quad word boundaries
    - Virtual memory trickier when datum spans 2 pages
- ▶ Compiler
  - Inserts gaps in structure to ensure correct alignment of fields

# Structures & Alignment

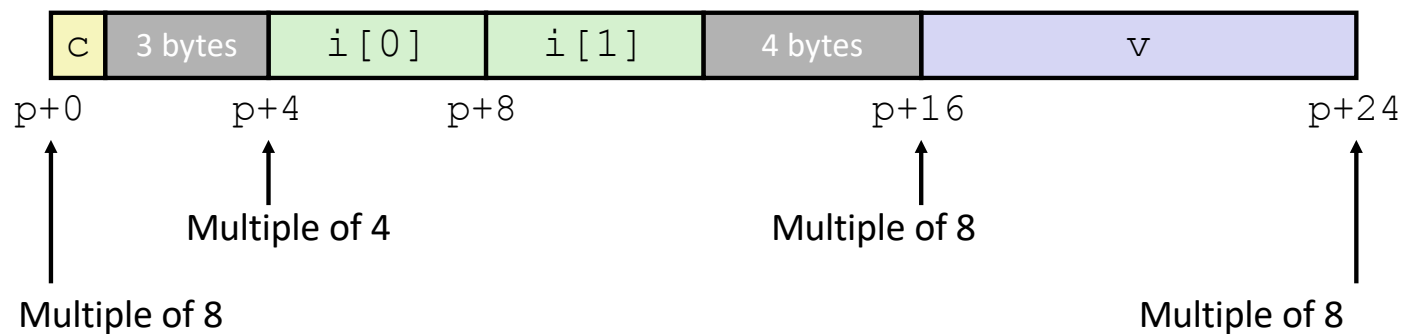
## ▶ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ▶ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



# Alignment Principles

- ▶ **Aligned Data**
  - Primitive data type requires  $K$  bytes
  - Address must be multiple of  $K$
  - Required on some machines; advised on x86-64
- ▶ **Motivation for Aligning Data**
  - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
    - Inefficient to load or store datum that spans quad word boundaries
    - Virtual memory trickier when datum spans 2 pages
- ▶ **Compiler**
  - Inserts gaps in structure to ensure correct alignment of fields

# Specific Cases of Alignment (x86-64)

- ▶ 1 byte: **char**, ...
  - no restrictions on address
- ▶ 2 bytes: **short**, ...
  - lowest 1 bit of address must be 0<sub>2</sub>
- ▶ 4 bytes: **int**, **float**, ...

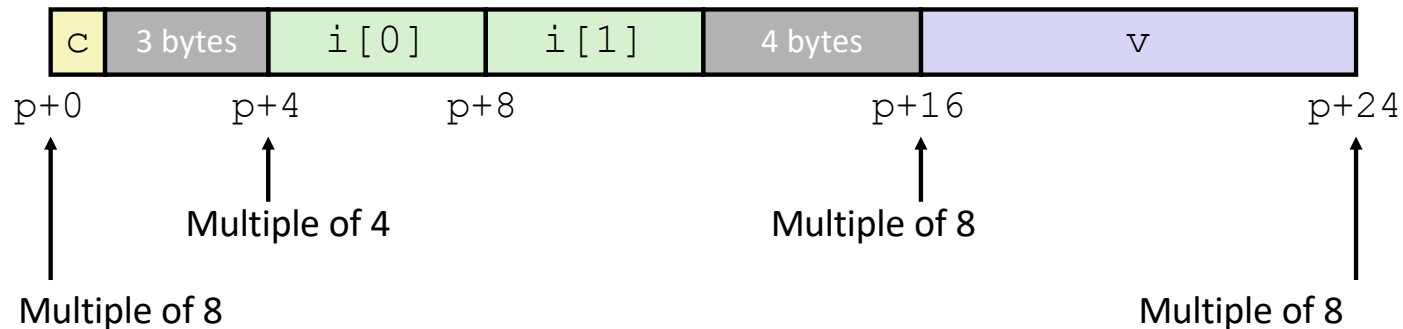
lowest 2 bits of address must be 00<sub>2</sub>
- ▶ 8 bytes: **double**, **long**, **char \***, ...
  - lowest 3 bits of address must be 000<sub>2</sub>
- ▶ 16 bytes: **long double** (GCC on Linux)
  - lowest 4 bits of address must be 0000<sub>2</sub>



# Satisfying Alignment with Structures

- ▶ Within structure:
  - Must satisfy each element's alignment requirement
- ▶ Overall structure placement
  - Each structure has alignment requirement  $K$ 
    - $K$  = Largest alignment of any element
  - Initial address & structure length must be multiples of  $K$
- ▶ Example:
  - $K = 8$ , due to **double** element

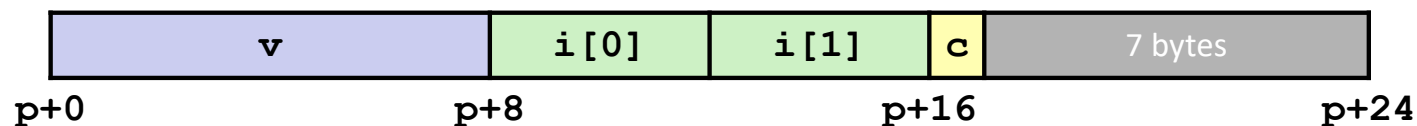
```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



# Meeting Overall Alignment Requirement

- ▶ For largest alignment requirement K
- ▶ Overall structure must be multiple of K

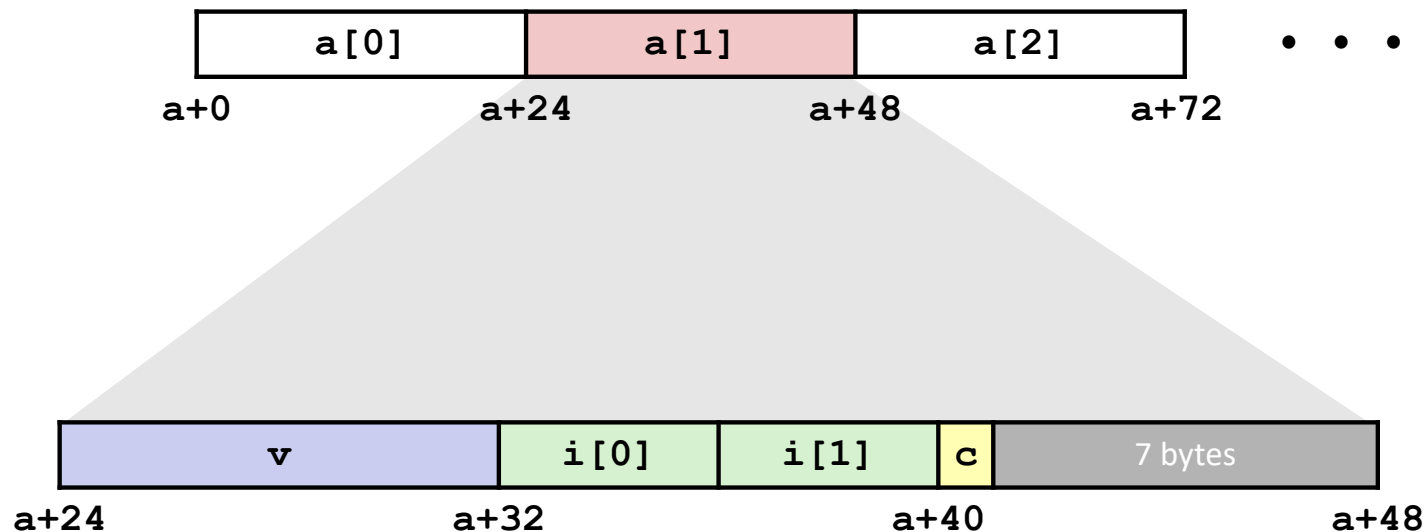
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



# Arrays of Structures

- ▶ Overall structure length multiple of K
- ▶ Satisfy alignment requirement for every element

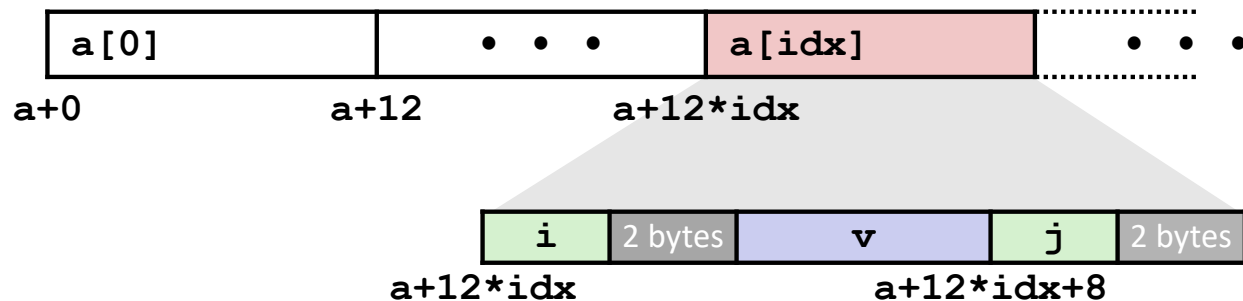
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Accessing Array Elements

- ▶ Compute array offset  $12 \cdot \text{idx}$ 
  - `sizeof(S3)`, including alignment spacers
- ▶ Element `j` is at offset 8 within structure
- ▶ Assembler gives offset `a+8`
  - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```

# Saving Space

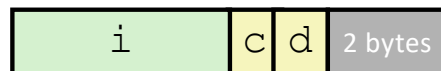
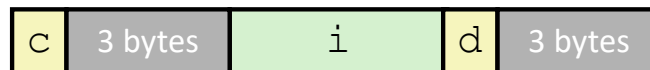
- ▶ Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- ▶ Effect (K=4)



# Rules

1. The whole structure must be multiple of  $k$ , where  $k$  is the size of largest primitive data size
2. Each member variables must be aligned according to their requirement